

Introducing Interrupts

Pr. Olivier Gruber

olivier.gruber@univ-grenoble-alpes.fr

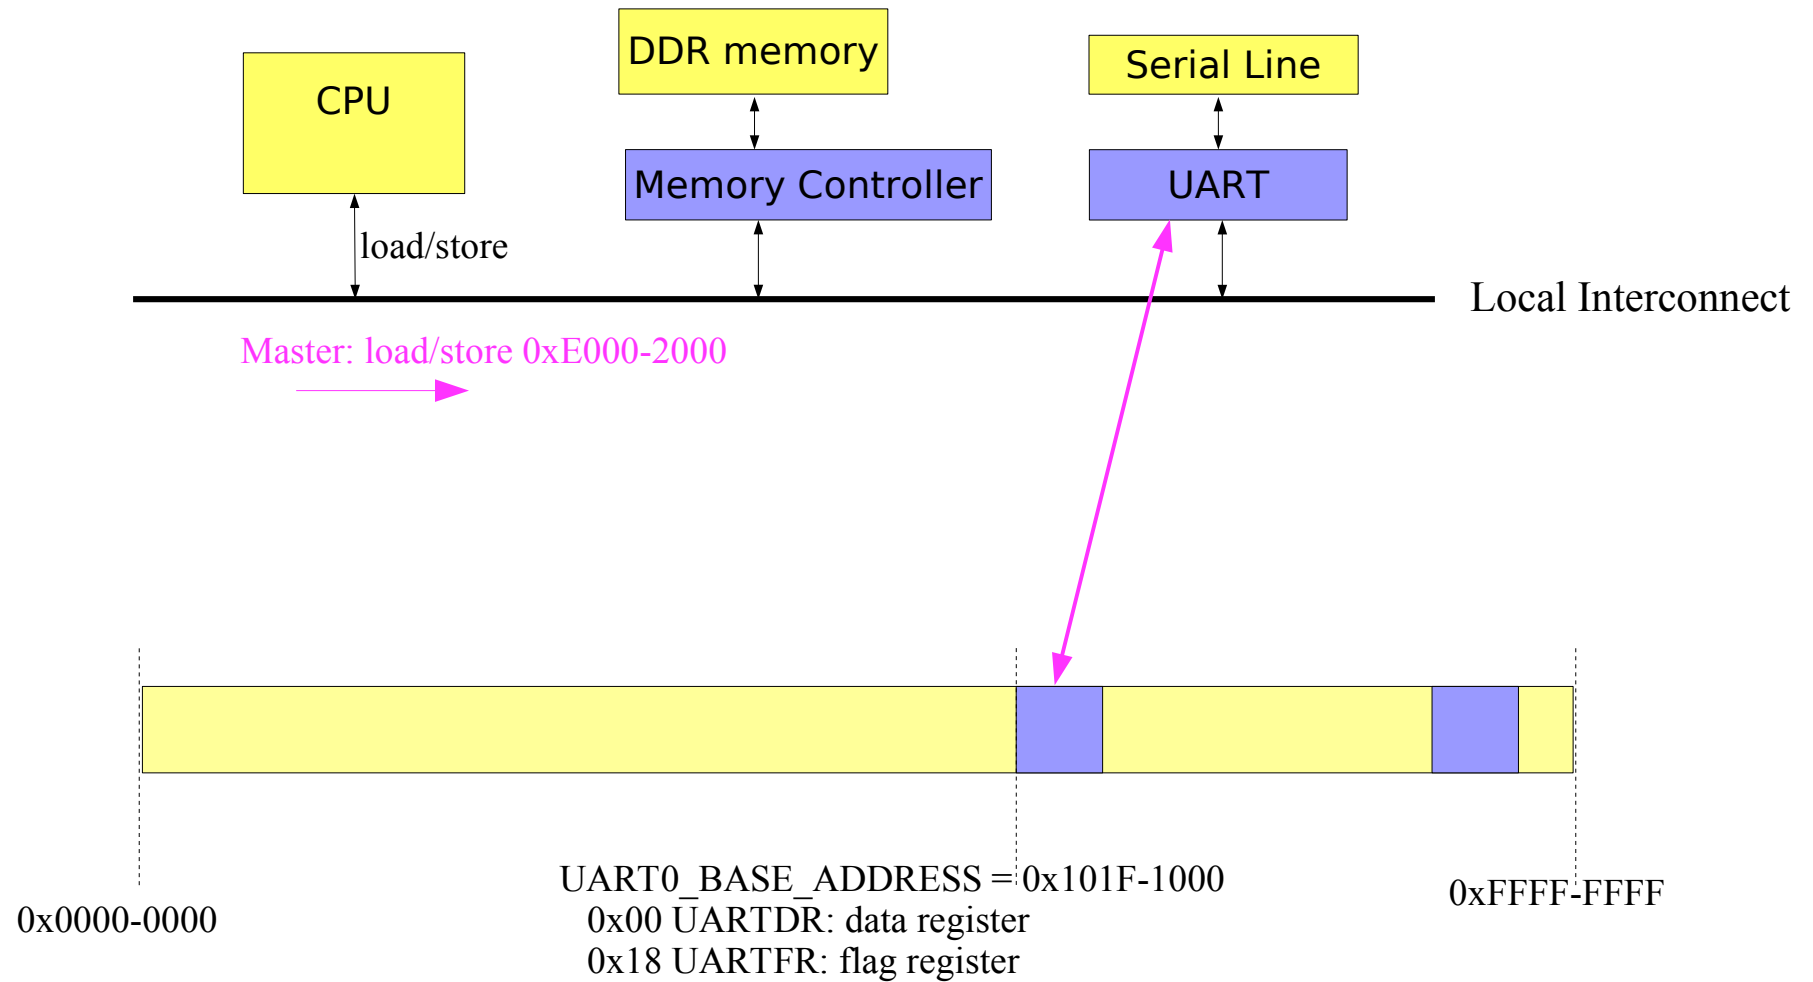
Laboratoire d'Informatique de Grenoble

Université de Grenoble-Alpes

- Quick recap
 - Hardware concepts
 - Basic examples
- Identify main coding challenges
 - Power consumption, timing constraints, race conditions
 - Better engineering of the code
- Discussing interrupts
 - Hardware perspective
 - Software perspective

Hardware MMIO Registers

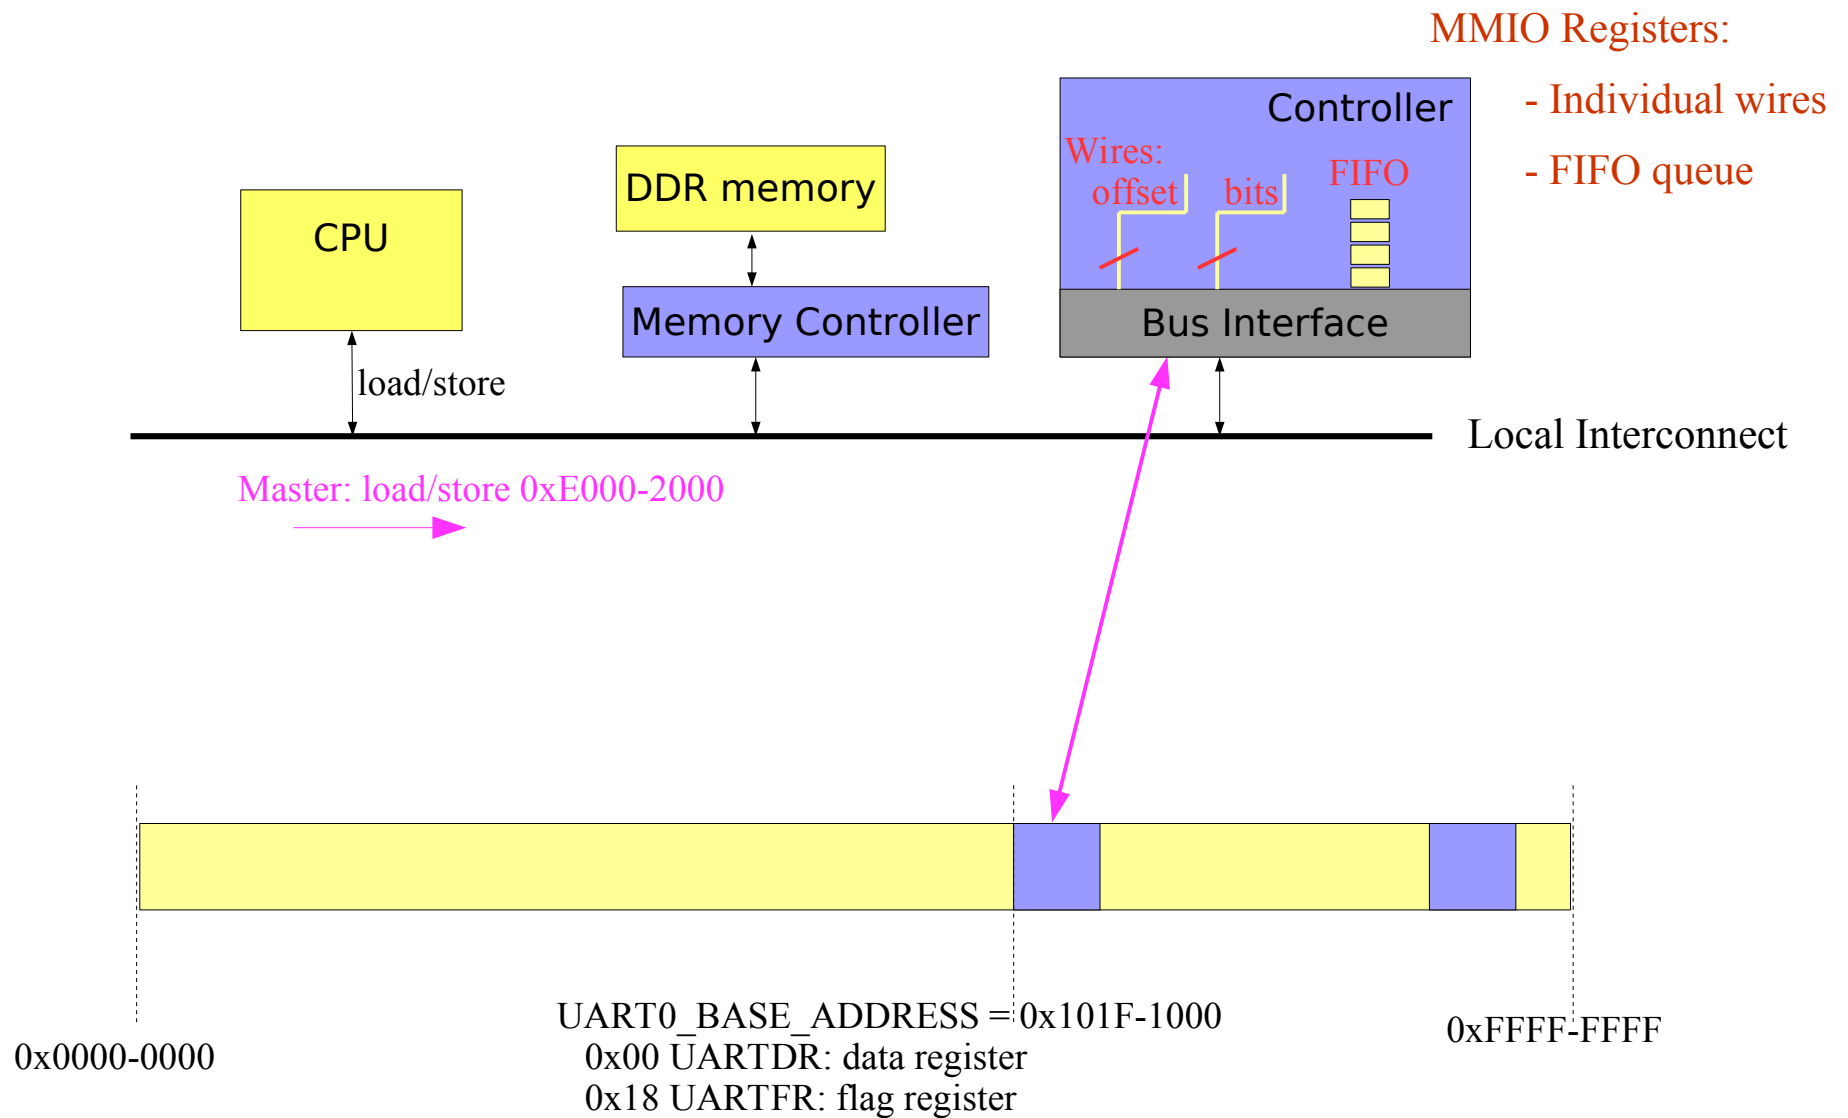
3



Memory-Mapped Input/Output Ranges

Hardware MMIO Registers

4



Bare Metal Programming – A simple example

5

- A button/led example
 - A forever loop
 - Check button status → set LED on or off accordingly

```
void c_entry() {  
  
    button_init_regs(BUTTON_MMIO_BASE_ADDRESS);  
    led_init_regs(LED_MMIO_BASE_ADDRESS);  
  
    for (;;) {  
        uint8_t status = button_get_status(BUTTON_MMIO_BASE_ADDRESS);  
        if (status & 0x01)  
            led_on(LED_MMIO_BASE_ADDRESS);  
        else  
            led_off(LED_MMIO_BASE_ADDRESS);  
    }  
}
```

Challenge?

Bare Metal Programming – A simple example

6

- A button/led example
 - A forever loop
 - Check button status → set LED on or off accordingly

```
void c_entry() {  
  
    button_init_regs(BUTTON_MMIO_BASE_ADDRESS);  
    led_init_regs(LED_MMIO_BASE_ADDRESS);  
  
    for (;;) {  
        uint8_t status = button_get_status(BUTTON_MMIO_BASE_ADDRESS);  
        if (status & 0x01)  
            led_on(LED_MMIO_BASE_ADDRESS);  
        else  
            led_off(LED_MMIO_BASE_ADDRESS);  
    }  
}
```

Challenge: continuous polling.

- processor never halts
- high power consumption
- short battery life

Bare Metal Programming – Another example

7

- A shell example via a serial line
 - Continuous polling of keyboard → high power consumption → short battery life
 - Other challenge such as timing constraints?

```
void c_entry() {  
    uart_init_regs(UART0);  
  
    unicode_t line[80];  
    int offset = 0;  
    for (;;) {  
        uint8_t code = uart_get_code(UART0);  
        if (code) {  
            // send the corresponding character to the screen  
            uart_put_byte(UART0, code);  
            if (code == '\n') { // do we have a full line?  
                interpret(line, offset);  
                offset = 0;  
            } else  
                line[offset++] = code;  
        }  
    }  
}
```

Bare Metal Programming – Another example

8

- A shell example via a serial line
 - Continuous polling of keyboard → high power consumption → short battery life
 - **Challenge**: not loosing characters because of long commands

Evaluation:

UART with 8-byte FIFO

115200 bps → over 14000 Bps

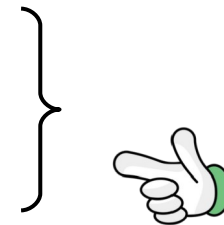
One byte every 70 us

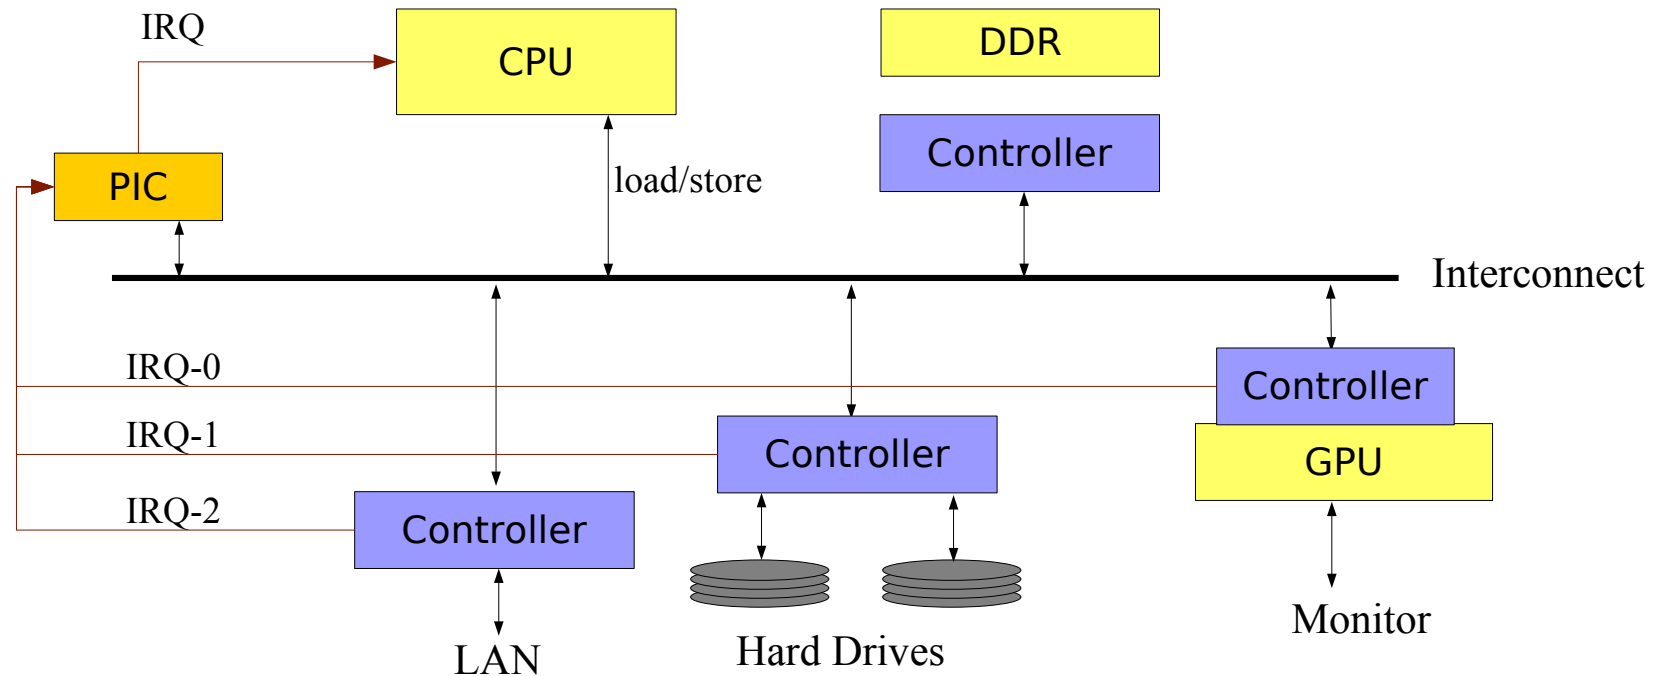
Dropping bytes after 560 us

560 us = (8 * 70us)

```
void c_entry() {  
  
    uart_init_regs(UART0);  
  
    unicode_t line[80];  
    int offset = 0;  
    for (;;) {  
        uint8_t code = uart_get_code(UART0);  
        if (code) {  
            // send the corresponding character to the screen  
            uart_put_byte(UART0, code);  
            if (code == '\n') { // do we have a full line?  
                interpret(line, offset);  
                offset = 0;  
            } else  
                line[offset++] = code;  
        }  
    }  
}
```


- Identified challenges
 - Continuous **polling**: waiting until a device has data available to read
 - Potential **spinning**: waiting until a device is ready to accept data
 - Timing constraints: limited time to pay attention to a device that requires attention
- Identified problems with our coding approach
 - Processor never halts, high power consumption, heat production
 - Potential loss of data if devices are not attended to in time
- Challenges
 - How do we allow the processor to halt?
 - **How do we know when a device needs attention?**





PIC: Programmable Interrupt Controller

A chipset to memorize and multiplex interrupt requests from hardware devices
Configured and controlled through mmio registers, like any other controller

- Device Controller Interrupts
 - Each device controller may raise one or more interrupts
 - Each interrupt is one dedicated electrical line to the PIC
- Programmable Interrupt Controller⁽¹⁾
 - A new device on the bus, configured through mmio registers
 - Multiplexer/demultiplexer for IRQs from devices
- New processor pin
 - **Interrupt signal** from the PIC
 - Like a door bell, asking the processor to pay attention to one or more devices
 - The door bell can be enabled or disabled by software

(1) May have different names such as GIC or VIC or APIC

- Interrupts/Traps are *exceptions* for the processor

The 7 known exceptions:

```
.globl _vector = 0x00000000
_vector:
    ldr    pc, _reset
    ldr    pc, _undefined_instruction
    ldr    pc, _software_interrupt
    ldr    pc, _prefetch_abort
    ldr    pc, _data_abort
    ldr    pc, _not_used
    ldr    pc, _irq
    ldr    pc, _fiq
```

↔ ldr pc, [pc+18]

The **irq_handler** is the next processing step for handling interrupts...

```
_reset:                .word undefined_instruction
_undefined_instruction: .word undefined_instruction
_software_interrupt:   .word software_interrupt
_prefetch_abort:      .word prefetch_abort
_data_abort:          .word data_abort
_not_used:             .word not_used
_irq:                .word _irq_handler
_fiq:                 .word fast_interrupt_service_routine

_irq_handler:
    ...
    bl irq_handler
    ...
```

ARM Example – Software perspective

13

Polling the PIC (via mmio reads)

Re-enabling interrupts at the PIC

```
void irq_handler(void) {  
    for (;;) {  
        int irq = pic_next_raised_irq();  
        switch (irq) {  
            case -1:  
                pic_enable_all_irqs();  
                return;  
            case UART0_IRQ:  
                uart0_isr();  
                break;  
        }  
    }  
}
```

The **interrupt service routine** for the UART0

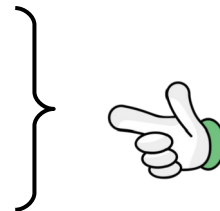
Suppose to figure out why the UART0 raised its interrupt....

Here we assume there are pending characters

For most device, **must not** forget to acknowledge the IRQ to the their device

```
void uart0_isr() {  
    ...  
    uart_irq_ack(UART0); // device-specific  
}
```















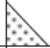
- PIC questions
 - How do we tell which interrupts has been raised?
 - How do we clear the interrupt status and why?
- UART questions
 - How do we tell which interrupts has been raised?
 - How do we clear the interrupt status and why?
- Hardware-software boundary questions
 - Processor mode when handling the interrupt
 - Assembly/C calling convention




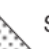



- ARM processors have different operating modes
 - **User**: normal program execution
 - **FIQ**: fast interrupt mode
 - **IRQ**: regular interrupt mode
 - **Supervisor**: protected mode for the operating system
 - **Abort**: when a page fault occurred when the MMU is turned on
 - **Undefined**: when an undefined instruction has encountered (supports emulation)
 - **System**: to run privileged operating system tasks
- User vs Privilege
 - The modes other than the User mode are considered privileged modes
 - They have full access to the processor and can reconfigure it

ARM Example – Processor Modes

16

Privileged modes						
Exception modes						
User	System	Supervisor	Abort	Undefined	Interrupt	Fast interrupt
R0	R0	R0	R0	R0	R0	R0
R1	R1	R1	R1	R1	R1	R1
R2	R2	R2	R2	R2	R2	R2
R3	R3	R3	R3	R3	R3	R3
R4	R4	R4	R4	R4	R4	R4
R5	R5	R5	R5	R5	R5	R5
R6	R6	R6	R6	R6	R6	R6
R7	R7	R7	R7	R7	R7	R7
R8	R8	R8	R8	R8	R8	 R8_fiq
R9	R9	R9	R9	R9	R9	 R9_fiq
R10	R10	R10	R10	R10	R10	 R10_fiq
R11	R11	R11	R11	R11	R11	 R11_fiq
R12	R12	R12	R12	R12	R12	 R12_fiq
R13	R13	 R13_svc	 R13_abt	 R13_und	 R13_irq	 R13_fiq
R14	R14	 R14_svc	 R14_abt	 R14_und	 R14_irq	 R14_fiq
PC	PC	PC	PC	PC	PC	PC

CPSR	CPSR	CPSR	CPSR	CPSR	CPSR	CPSR
		 SPSR_svc	 SPSR_abt	 SPSR_und	 SPSR_irq	 SPSR_fiq

 indicates that the normal register used by User or System mode has been replaced by an alternative register specific to the exception mode

ARM Example – Interrupt Mode

17

When in interrupt mode → interrupts are disabled on the processor

Note that the interrupt mode has its own **stack**, **link**, and **SPSR** registers.

R0 – R12: general-purpose registers

R13: **sp** (stack pointer)

R14: **lr** (link register)

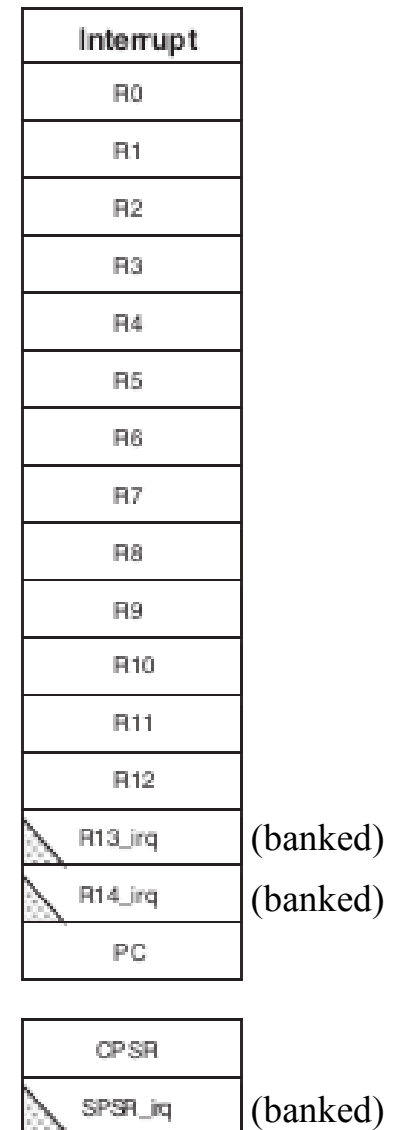
R15: **pc** (program counter)

CPSR: Current Program Status Register

SPSR: Saved Program Status Register

The Current Program Status Register (CPSR) is accessible in all processor modes. It contains condition code flags, interrupt disable bits, the current processor mode, and other status and control information.

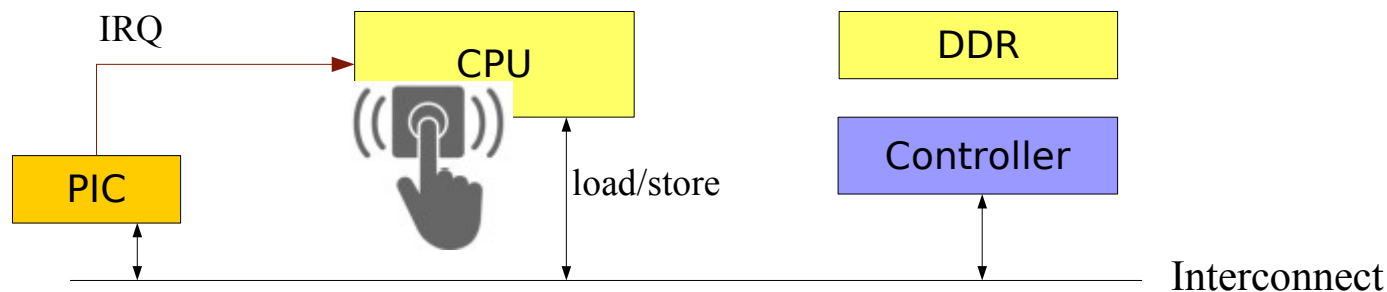
Each exception mode also has a Saved Program Status Register (SPSR), that is used to preserve the value of the CPSR when the associated exception occurs.



Startup initialization:

- Setting up the IRQ-mode stack
(define in the linker script as another stack)
- Enabling the IRQ interrupts at the processor⁽¹⁾

```
.global _setup_irqs
_enable_irq:
    /* get Program Status Register */
    MRS r0, cpsr
    /* go in IRQ mode */
    BIC r1, r0, #0x1F
    ORR r1, r1, #0x12
    MSR cpsr, r1
    /* set IRQ stack */
    LDR sp, =irq_stack_top
    /* Enable IRQs */
    BIC r0, r0, #0x80
    /* go back in Supervisor mode */
    MSR cpsr, r0
    mov pc, lr
```



(1) Meaning the PIC can ring the door on processor and asking the processor to handle an interrupt.
The processor will do it, eventually, but not immediately.

IRQ handler:

- Fix the link register
- Save all non-banked registers and link register
- Upcall the C function
- Restore all non-banked registers and the pc

```
_irq_handler:
    sub lr,lr,#4
    stmfd sp!, {r0-r12,lr}
    bl irq_handler
    ldmfd sp!, {r0-r12,pc}^

void irq_handler() {

}
```

Note:

The ^ symbol at the end of the LDMFD instruction means that the CPSR will be restored from the save SPSR. This only happens if the pc is loaded by the LDMFD instruction.

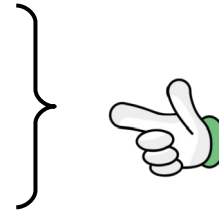
Restoring the SPSR means re-enabling the interrupts on the processor. Indeed, if we are processing an interrupt, they were enabled on the processor. By going in interrupt mode, the processor disabled further interrupts. By returning from interrupt mode, the processor therefore restores the ability to handle further interrupts.

- PIC questions

- How do we tell which interrupts has been raised?
- How do we clear the interrupt status and why?

- UART questions

- How do we tell which interrupts has been raised?
- How do we clear the interrupt status and why?



- ISA-related questions

- Processor mode when handling the interrupt
- Assembly/C calling convention

- How do we tell which interrupts has been raised?
 - The UARTIMSC Register is the interrupt mask **set/clear**⁽²⁾ register. It is a read/write register.
Masked interrupts are enabled
 - The UARTRIS Register is the raw interrupt status register. It is a read-only register that returns the current raw status value, prior to masking
 - The UARTMIS Register is the masked interrupt status register. It is a read-only register. This register returns the current masked status value of the corresponding interrupt. A write has no effect
- How do we clear the interrupt status at the UART?
 - The UARTICR Register is the interrupt clear register and is write-only. On a write of 1, the corresponding interrupt is cleared. A write of 0 has no effect.
 - **Note:** the **receive interrupt** is also cleared by performing at least one read of the receive FIFO, it may be cleared by clearing the interrupt in UARTICR register.

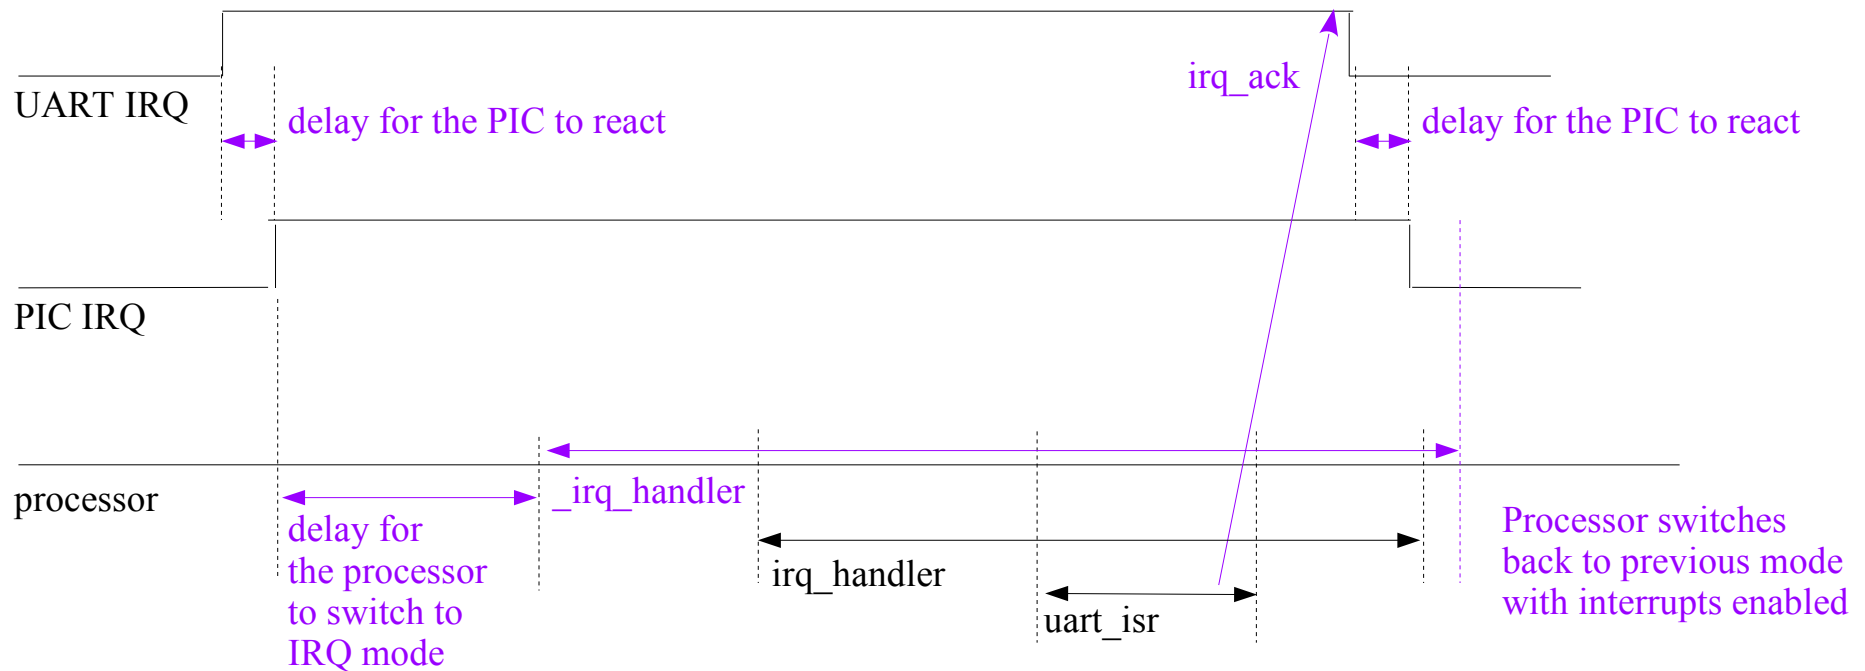
(1) PrimeCell UART (PL011) Technical Reference Manual (DDIO183G)

(2) Set/Clear register: write 1s to set the corresponding mask of that interrupt, write 0s to clear it

UART Management – PL011 Example

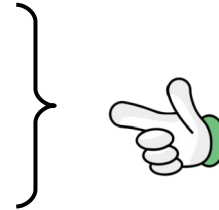
22

- Why do we need to clear the interrupt at the UART (PL011)?
 - How would the UART otherwise know when the interrupt has been handled by software?
 - It needs to know in order to know when it can raise another interrupt if necessary



- PIC questions

- How do we tell which interrupts has been raised?
- How do we clear the interrupt status and why?



- UART questions

- How do we tell which interrupts has been raised?
- How do we clear the interrupt status and why?

- ISA-related questions

- Processor mode when handling the interrupt
- Assembly/C calling convention

- How do we tell which interrupts has been raised?
 - The read-only VIC_IRQ_STATUS⁽²⁾ register provides the status of interrupts [31:0] after IRQ masking. **Masked interrupts are enabled.**
 - The read/write VIC_INT_ENABLE⁽²⁾ register enables the interrupt request lines, by masking the interrupt sources for the IRQ interrupt.
 - The write-only VIC_INT_ENCLEAR⁽²⁾ register clears bits in the VIC_INT_ENABLE register.
- How do we clear the interrupt status at the VIC?
 - Not necessary for the PL190 because interrupt inputs must be level sensitive, active HIGH, and held asserted until the interrupt service routine clears the interrupt.
 - For PIC that handles edge-triggered interrupts, an ACK is usually necessary for the same reason as for the UART, the PIC needs to know when a particular interrupt has been handled by software

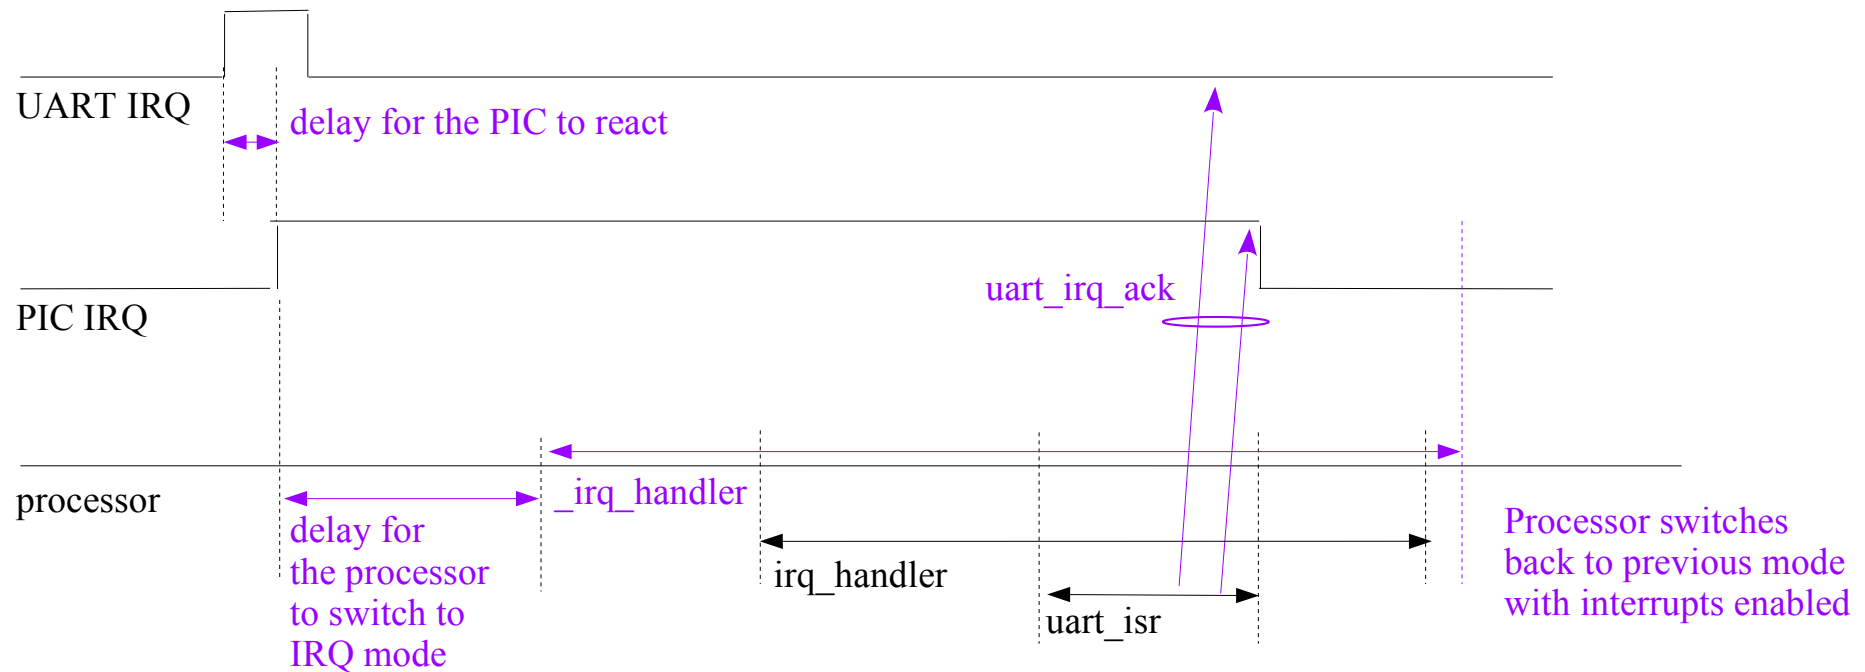
(1) PrimeCell Vectored Interrupt Controller (PL190) Technical Reference Manual

(2) Named in the doc VICIRQSTATUS, VICINTENABLE, VICINTENCLEAR

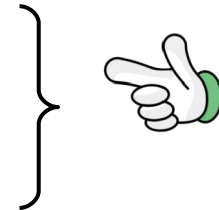
PIC Management – Edge-Triggered

25

Example of a PIC that handles edge-triggered IRQs from controllers and allowing individual acknowledgment of such IRQs, (typically the case on Intel boards and the Advanced Programmable Interrupt Controller (APIC))



- Identified challenges
 - Continuous **polling**: waiting until a device has data available to read
 - Potential **spinning**: waiting until a device is ready to accept data
 - Timing constraints: limited time to pay attention to a device that requires attention
- Identified problems with our coding approach
 - Processor never halts, high power consumption, heat production
 - Potential loss of data if devices are not attended to in time
- Challenges
 - **How do we allow the processor to halt?**
 - How do we know when a device needs attention?



Interrupts – Software perspective

27

- New instruction: **halt**⁽¹⁾
 - Enables interrupts
 - Halts only if the processor interrupt signal is inactive
 - Wakes up when the processor interrupt signal becomes active

```
void uart0_isr() {  
    ...  
    uart_irq_ack(UART0);  
}  
  
void main() {  
    uart_init_regs(UART0);  
    for (;;) {  
        ...  
        wfi();  
    }  
}
```

Don't forget to halt the processor
when there is nothing more to do...
Here, we *Wait For Interrupts* (wfi)

(1): Halt instruction has different names for different processors

May not be an instruction on some processor, but a flag in a control register:

ARMv5:	MCR
ARMv7:	wfi
Intel:	hlt

Interrupts – Software perspective

28

- New instruction: **halt**
 - Enables interrupts
 - Halts only if the processor interrupt signal is inactive
 - Wakes up when the processor interrupt signal becomes active

```
void uart0_isr() {  
    ...  
    uart_irq_ack(UART0);  
}  
  
void main() {  
    uart_init_regs(UART0);  
    for (;;) {  
        ...  
        wfi();  
    }  
}
```

Missing something, are we not?

We need a communication channel
between the ISR and the main loop...

Right?

- The ISR gathers the characters...
- The main processes the characters...

Our Shell Example with IRQs

29

```
uint8_t line[80];
int offset = 0;
int full_line = 0;

void uart0_isr() {
    uint8_t code = uart_get_byte(UART0);
    while (code) {
        if (code == '\n') { // full line?
            full_line = 1;
            break;
        } else
            line[offset++] = code;
        code = uart_get_byte(UART0);
    }
    uart_irq_ack(UART0);
}

void main() {
    uart_init_regs(UART0);
    for (;;) {
        if (full_line) {
            interpret(line, offset);
            offset = 0;
            full_line = 0;
        }
        halt();
    }
}
```

Looks good... right?

We introduced a buffer to gather the characters from the UART, one line at a time.

Our Shell Example with IRQs

30

```
uint8_t line[80];
int offset = 0;
int full_line = 0;

void uart0_isr() {
    uint8_t code = uart_get_byte(UART0);
    while (code) {
        if (code == '\n') { // full line?
            full_line = 1;
            break;
        } else
            line[offset++] = code;
        code = uart_get_byte(UART0);
    }
    uart_irq_ack(UART0);
}

void main() {
    uart_init_regs(UART0);
    for (;;) {
        if (full_line) {
            interpret(line, offset);
            offset = 0;
            full_line = 0;
        }
        halt();
    }
}
```

Not really...

Even though it may work for a while...

It may run, but it is not correct!

Do you see what is the problem?

Our Shell Example with IRQs

31

```
uint8_t line[80];
int offset = 0;
int full_line = 0;

void uart0_isr() {
    uint8_t code = uart_get_byte(UART0);
    while (code) {
        if (code == '\n') { // full line?
            full_line = 1;
            break;
        } else
            line[offset++] = code;
        code = uart_get_byte(UART0);
    }
    uart_irq_ack(UART0);
}

void main() {
    uart_init_regs(UART0);
    for (;;) {
        if (full_line) {
            interpret(line, offset);
            offset = 0;
            full_line = 0;
        }
        halt();
    }
}
```

Race condition between:

- (a) UART ISR
- (b) Line parsing to interpret

over the shared line buffer!

Our Shell Example with IRQs – First Solution

32

```
uint8_t line[80];
int offset = 0;
int full_line = 0;

void uart0_isr() {
    uint8_t code = uart_get_byte(UART0);
    while (code) {
        if (code == '\n') { // full line?
            full_line = 1;
            pic_disable_irqs();
            break;
        } else {
            line[offset++] = code;
            code = uart_get_byte(UART0);
        }
    }
    uart_irq_ack(UART0);
}

void main() {
    uart_init_regs(UART0);
    for (;;) {
        if (full_line) {
            interpret(line, offset);
            offset = 0;
            full_line = 0;
            pic_enable_irqs();
        }
        halt();
    }
}
```

First solution → disable/enable IRQs

- Disable IRQs when a full line has been captured,
- Re-enable IRQs when the full line has been interpreted.

Challenge?

Our Shell Example with IRQs – Second Solution

33

```
#define MAX_CHARS 256
volatile uint8_t buffer[MAX_CHARS];
volatile int head = 0;
volatile int tail = 0;

void uart0_isr() {
    uint8_t code = uart_get_byte(UART0);
    while (code) {
        int next = (head + 1) % MAX_CHARS;
        if (next==tail) return;
        buffer[head] = code;
        head = next;
        code = uart_get_byte(UART0);
    }
}

void main() {
    int eol = 0;
    uart_init_regs(UART0);
    for (;;) {
        while (eol!=head) {
            eol = (eol + 1) % MAX_CHARS;
            if (buffer[eol] == '\n') {
                interpret(buffer,tail,eol);
                eol = (eol + 1) % MAX_CHARS;
                tail = eol;
            }
        }
        halt();
    }
}
```

Second solution → circular buffer

Note: a lock-free design is necessary

Why? Because blocking the ISR would in fact deadlock the execution!

Do you see why?

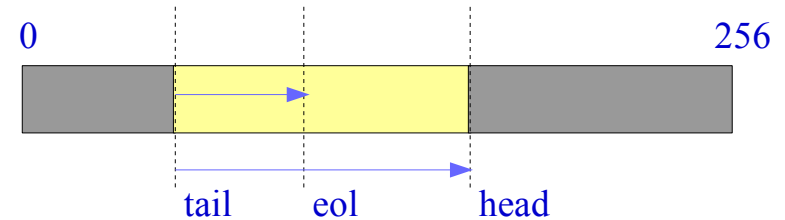
Using a Lock-free Circular Buffer

34

```
#define MAX_CHARS 256
volatile uint8_t buffer[MAX_CHARS];
volatile int head = 0;
volatile int tail = 0;

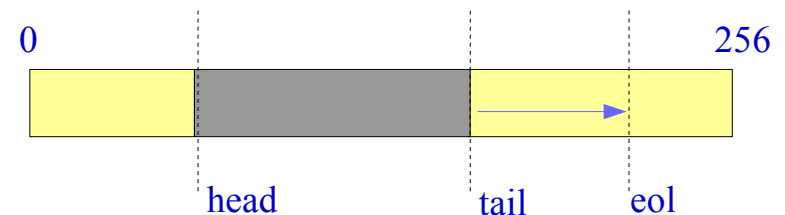
void uart0_isr() {
    uint8_t code = uart_get_byte(UART0);
    while (code) {
        int next = (head + 1) % MAX_CHARS;
        if (next==tail) return;
        buffer[next] = code;
        head = next;
        code = uart_get_byte(UART0);
    }
}

void main() {
    int eol = 0;
    uart_init_regs(UART0);
    for (;;) {
        while (eol!=head) {
            eol = (eol + 1) % MAX_CHARS;
            if (buffer[eol] == '\n') {
                interpret(buffer, tail, eol);
                eol = (eol + 1) % MAX_CHARS;
                tail = eol;
            }
        }
        halt();
    }
}
```



Circular Buffer:

Acts as a lock-free communication channel between the ISR and the main loop



Using a Lock-free Circular Buffer

35

```
#define MAX_CHARS 256
volatile uint8_t buffer[MAX_CHARS];
volatile int head = 0;
volatile int tail = 0;

void uart0_isr() {
    uint8_t code = uart_get_byte(UART0);
    while (code) {
        int next = (head + 1) % MAX_CHARS;
        if (next==tail) return;
        buffer[head] = code;
        head = next;
        code = uart_get_byte(UART0);
    }
}

void main() {
    int eol = 0;
    uart_init_regs(UART0);
    for (;;) {
        while (eol!=head) {
            eol = (eol + 1) % MAX_CHARS;
            if (buffer[eol] == '\n') {
                interpret(buffer,tail,eol);
                eol = (eol + 1) % MAX_CHARS;
                tail = eol;
            }
        }
        halt();
    }
}
```

Circular Buffer

Warning: the circular buffer may be full...

Can you lose characters then? Of course...

This means that the main loop must be able to keep up with incoming characters before the circular buffer fills up

There is no free lunch!

The circular buffer only buys your code more time to react...

- We strongly suggest the following
 - Clearly separate the core exception and IRQ support
 - Exception vector
 - Global interrupt handler
 - Clearly organize the UART code (uart.c and uart.h)
 - Including the UART interrupt service routine
 - Circular buffer
- Cleaner code / easier development
 - Easier to read
 - Easier to understand
 - Easier to debug
 - Easier to evolve

- Architecture Reference Manual
 - For the processor, the VersatilePB uses a ARMv5 processor
- Technical Reference Manuals
 - PrimeCell Vectored Interrupt Controller (PL190)
 - PrimeCell UART (PL011)
- How to read them?
 - (1) do not get scared, do not run away to Google land...
 - Understand they target two audiences: hardware integrators and software developers
 - So focus on the parts that are for software developers
 - (2) understand the typical outline
 - Functional Overview ← more hardware than software
 - Programmers Model ← more software than hardware